

Twitter Sentimental Analysis

AUTHOR: MOHANKRISHNA

Importing Libraries

```
In [3]: # Importing necessary libraries
import re
import pickle
import numpy as np
import pandas as pd

import seaborn as sns
from wordcloud import WordCloud
import matplotlib.pyplot as plt

# Setting a more appealing style for plots
sns.set(style='whitegrid')
plt.style.use('ggplot') # Use ggplot style for better visual appeal
```

Importing Dataset

```
In [5]: # Importing the dataset and processing it
DATASET_COLUMNS = ["sentiment", "ids", "date", "flag", "user", "text"]
DATASET_ENCODING = "ISO-8859-1"

# Reading the dataset and selecting only necessary columns
dataset = pd.read_csv('twitter_sentiment_data.csv', encoding=DATASET_ENCODING,

# Renaming the columns to 'sentiment' and 'text'
dataset.columns = ['sentiment', 'text']

# Replacing sentiment values for easier interpretation (0 = negative, 1 = positive)
dataset['sentiment'] = dataset['sentiment'].replace(4, 1)

# Dropping any missing values for cleanliness
dataset.dropna(inplace=True)

# Displaying the first few rows of the dataset for inspection
print(dataset.head())
```

	sentiment	text
0	0	is upset that he can't update his Facebook by ...
1	0	@Kenichan I dived many times for the ball. Man...
2	0	my whole body feels itchy and like its on fire
3	0	@nationwideclass no, it's not behaving at all....
4	0	@Kwesidei not the whole crew

Text Preprocessing is traditionally an important step for Natural Language Processing (NLP) tasks. It transforms text into a more digestible form so that machine learning algorithms can perform better.

The Preprocessing steps taken are:

1. Lower Casing: Each text is converted to lowercase. Replacing URLs: Links starting with "http" or "https" or "www" are replaced by "URL".
2. Replacing Emojis: Replace emojis by using a pre-defined dictionary containing emojis along with their meaning. (eg: ":)" to "EMOJIsmile")
3. Replacing Usernames: Replace @Usernames with word "USER". (eg: "@Kaggle" to "USER")
4. Removing Non-Alphabets: Replacing characters except Digits and Alphabets with a space.
5. Removing Consecutive letters: 3 or more consecutive letters are replaced by 2 letters. (eg: "Heyyyy" to "Heyy")
6. Removing Short Words: Words with length less than 2 are removed.
7. Removing Stopwords: Stopwords are the English words which does not add much meaning to a sentence. They can safely be ignored without sacrificing the meaning of the sentence. (eg: "the", "he", "have")
8. Lemmatizing: Lemmatization is the process of converting a word to its base form. (e.g: "Great" to "Good")

```
In [7]: # Defining dictionary containing all emojis with their meanings.
emojis = {' :)': 'smile', ' :-)': 'smile', ' ;d': 'wink', ' :-E': 'vampire', ' :( ':
        'sad', ' :-<': 'sad', ' :P': 'raspberry', ' :O': 'surprised',
        ' :-@': 'shocked', ' :@': 'shocked', ' :-$: 'confused', ' :\\': 'annoyed',
        ' :#': 'mute', ' :X': 'mute', ' :^)': 'smile', ' :-&': 'confused', ' $ _ $':
        '@@': 'eyeroll', ' :-!': 'confused', ' :-D': 'smile', ' :-0': 'yell', '
        <(-_-)>': 'robot', ' d[ _ _ ]b': 'dj', ' :'-)': 'sadsmile', ' ;)': 'wink',
        ' ;-)': 'wink', ' O:-)': 'angel', ' O*-)': 'angel', ' (: -D': 'gossip', ' =^

## Defining set containing all stopwords in english.
stopwordlist = ['a', 'about', 'above', 'after', 'again', 'ain', 'all', 'am', 'an',
                'and', 'any', 'are', 'as', 'at', 'be', 'because', 'been', 'before',
                'being', 'below', 'between', 'both', 'by', 'can', 'd', 'did', 'do',
                'does', 'doing', 'down', 'during', 'each', 'few', 'for', 'from',
                'further', 'had', 'has', 'have', 'having', 'he', 'her', 'here',
                'hers', 'herself', 'him', 'himself', 'his', 'how', 'i', 'if', 'in',
                'into', 'is', 'it', 'its', 'itself', 'just', 'll', 'm', 'ma',
                'me', 'more', 'most', 'my', 'myself', 'now', 'o', 'of', 'on', 'once',
                'only', 'or', 'other', 'our', 'ours', 'ourselves', 'out', 'own', 's',
                's', 'same', 'she', 'shes', 'should', 'shouldve', 'so', 'some', 'st',
                't', 'than', 'that', 'thatll', 'the', 'their', 'theirs', 'them',
                'themselves', 'then', 'there', 'these', 'they', 'this', 'those',
                'through', 'to', 'too', 'under', 'until', 'up', 've', 'very', 'was',
                'we', 'were', 'what', 'when', 'where', 'which', 'while', 'who', 'who',
                'why', 'will', 'with', 'won', 'y', 'you', 'youd', 'youll', 'youre',
                'youve', 'your', 'yours', 'yourself', 'yourselves']
```

```
In [8]: from nltk.stem import WordNetLemmatizer
```

```
In [9]: def preprocess(textdata):
    processedText = []

    # Create Lemmatizer and Stemmer
    wordLemm = WordNetLemmatizer()

    # Regex patterns
    urlPattern = r"((http://)[^ ]*|(https://)[^ ]*|( www\.)[^ ]*)"
    userPattern = '@[^\s]+'
    alphaPattern = "[^a-zA-Z0-9]"
    sequencePattern = r"(\.)\1\1+"
    seqReplacePattern = r"\1\1"

    for tweet in textdata:
        tweet = tweet.lower()

        # Replace all URLs with 'URL'
        tweet = re.sub(urlPattern, ' URL', tweet)
        # Replace all emojis.
        for emoji in emojis.keys():
            tweet = tweet.replace(emoji, "EMOJI" + emojis[emoji])
        # Replace @USERNAME to 'USER'.
        tweet = re.sub(userPattern, ' USER', tweet)
        # Replace all non alphabets.
        tweet = re.sub(alphaPattern, " ", tweet)
        # Replace 3 or more consecutive letters by 2 letter.
        tweet = re.sub(sequencePattern, seqReplacePattern, tweet)
        tweetwords = ''
        for word in tweet.split():
            # Checking if the word is a stopword.
            # If word not in stopwordlist
            if len(word) > 1:

                word = wordLemm.lemmatize(word)

                tweetwords += (word + ' ')
        processedText.append(tweetwords)

    return processedText
```

```
In [10]: # Assuming 'dataset' has already been loaded and contains the 'text' column

import time

# Function for text preprocessing (example function, adjust as needed)
def preprocess(texts):
    # Example preprocessing: converting to lowercase and removing special characters
    return [re.sub(r'\W', ' ', str(text).lower()) for text in texts]

# Record the start time
t = time.time()

# Preprocess the 'text' column from the dataset
processed_text = preprocess(dataset['text'])

print(f'Text Preprocessing complete.')
print(f'Time Taken: {round(time.time()-t)} seconds')

# Display the first few rows of the processed text
print(processed_text[:5])
```

Text Preprocessing complete.

Time Taken: 17 seconds

```
['is upset that he can t update his facebook by texting it and might cry a
s a result school today also blah ', ' kenichan i dived many times for the
ball managed to save 50 the rest go out of bounds', 'my whole body feels i
tchy and like its on fire ', ' nationwideclass no it s not behaving at all
i m mad why am i here because i can t see you all over there ', ' kwesidei
not the whole crew ']
```

Analysing the data

Now we're going to analyse the preprocessed data to get an understanding of it. We'll plot **Word Clouds** for **Positive and Negative** tweets from our dataset and see which words occur the most.

Word-Cloud for Negative tweets

[illegible]

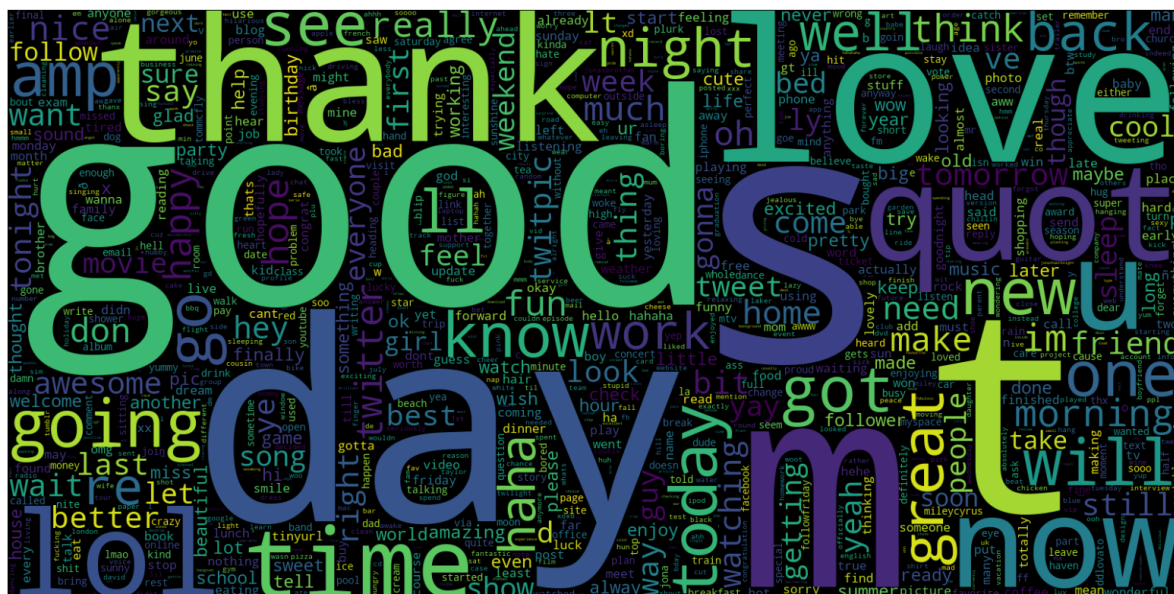
Word-Cloud for Positive tweets

```
In [15]: # Assuming 'processed_text' was defined after the preprocessing step

# Extract the positive sentiment data (assuming it's from the 800,000th entry of processed_text)
data_pos = processed_text[800000:] # Correct variable name

# Create a WordCloud for the positive text data
wc = WordCloud(max_words=1000, width=1600, height=800, collocations=False).generate(data_pos)

# Display the WordCloud
plt.figure(figsize=(20, 20))
plt.imshow(wc, interpolation='bilinear')
plt.axis('off') # Turn off the axis for better visualization
plt.show()
```



Splitting the data

```
In [17]: from sklearn.model_selection import train_test_split

In [18]: from sklearn.model_selection import train_test_split

# Ensure that 'processed_text' and 'dataset["sentiment"]' are defined
# processed_text contains preprocessed tweet text
# sentiment contains the target variable (Labels)

X_train, X_test, y_train, y_test = train_test_split(processed_text, dataset['sentiment'],
                                                    test_size=0.05, random_state=42)

print('Data Split done.')

Data Split done.
```

```
In [19]: from sklearn.feature_extraction.text import TfidfVectorizer
```

```
In [20]: from sklearn.feature_extraction.text import TfidfVectorizer

# Create and fit the TfidfVectorizer
vectorizer = TfidfVectorizer(ngram_range=(1,2), max_features=50000)
vectorizer.fit(X_train)

print('Vectorizer fitted')
# Use get_feature_names_out() instead of get_feature_names()
print('No. of feature_words: ', len(vectorizer.get_feature_names_out()))
```

```
Vectorizer fitted
No. of feature_words: 50000
```

```
In [21]: X_train = vectorizer.transform(X_train)
X_test = vectorizer.transform(X_test)
print(f'Data Transformed')
```

```
Data Transformed
```

Creating and Evaluating Models

Creating 3 different types of model of our sentimental analysis problems.

- **Bernoulli Naive Baye(Bernoulli)**
- **Linear Support Vector Classificatio (LinearSVC)**
- **Logistic Regression (LR)**

```
In [23]: from sklearn.svm import LinearSVC
from sklearn.naive_bayes import BernoulliNB
from sklearn.linear_model import LogisticRegression
```

Evaluation Model Function

```
In [25]: from sklearn.metrics import confusion_matrix, classification_report
```



```
In [26]: # Importing necessary Libraries
import re
import pickle
import numpy as np
import pandas as pd

import seaborn as sns
from wordcloud import WordCloud
import matplotlib.pyplot as plt

# Setting a more appealing style for plots
sns.set(style='whitegrid')
plt.style.use('ggplot') # Use ggplot style for better visual appeal
```

```
In [27]: def model_evaluate(model):
    y_pred = model.predict(X_test)

    # classification report
    print(classification_report(y_test, y_pred))

    # confusion report
    cf_matrix = confusion_matrix(y_test, y_pred)

    categories = ['Negative', 'Positive']

    group_names = ['True Neg', 'False Pos', 'False Neg', 'True Pos']

    group_percentages = ['{0:.2%}'.format(value) for value in cf_matrix.flatten() if value > 0]

    labels = [f'{v1}\n{v2}' for v1, v2 in zip(group_names, group_percentages)]
    labels = np.asarray(labels).reshape(2,2)


    sns.heatmap(cf_matrix, annot = labels, cmap = 'Blues', fmt = '',
                xticklabels = categories, yticklabels = categories)

    plt.xlabel("Predicted values", fontdict = {'size':14}, labelpad = 10)
    plt.ylabel("Actual values", fontdict = {'size':14}, labelpad = 10)
    plt.title ("Confusion Matrix", fontdict = {'size':18}, pad = 20)
```


BernoulliNB Model

```
In [29]: BNBmodel = BernoulliNB(alpha = 2)
BNBmodel.fit(X_train, y_train)
model_evaluate(BNBmodel)
```

	precision	recall	f1-score	support
0	0.80	0.78	0.79	39986
1	0.78	0.80	0.79	40014
accuracy			0.79	80000
macro avg	0.79	0.79	0.79	80000
weighted avg	0.79	0.79	0.79	80000



LinearSVC Model

```
In [31]: SVCmodel = LinearSVC()
SVCmodel.fit(X_train, y_train)
model_evaluate(SVCmodel)
```

C:\Users\mohan\anaconda3\Lib\site-packages\sklearn\svm_classes.py:31: Future Warning: The default value of `dual` will change from `True` to `'auto'` in 1.5. Set the value of `dual` explicitly to suppress the warning.
warnings.warn(

	precision	recall	f1-score	support
0	0.83	0.81	0.82	39986
1	0.81	0.83	0.82	40014
accuracy			0.82	80000
macro avg	0.82	0.82	0.82	80000
weighted avg	0.82	0.82	0.82	80000


Confusion Matrix



Logistic Regression Model

```
In [33]: LRmodel = LogisticRegression(C =2, max_iter=1000, n_jobs=1)
LRmodel.fit(X_train, y_train)
model_evaluate(LRmodel)
```

	precision	recall	f1-score	support
0	0.82	0.81	0.82	39986
1	0.81	0.83	0.82	40014
accuracy			0.82	80000
macro avg	0.82	0.82	0.82	80000
weighted avg	0.82	0.82	0.82	80000



We can clearly see that the **Logistic Regression Model** performs the best out of all the different models that we tried. It achieves nearly 82% accuracy while classifying the sentiment of a tweet.

Although it should also be noted that the BernoulliNB Model is the fastest to train and predict on. It also achieves 80% accuracy while calssifying.

Saving the model

```
In [36]: file = open('vectoriser-ngram-(1,2).pickle','wb')
pickle.dump(vectorizer, file)
file.close()

file = open('Sentiment-LR.pickle','wb')
pickle.dump(LRmodel, file)
file.close()

file = open('Sentiment-BNB.pickle','wb')
pickle.dump(BNBmodel, file)
file.close()
```

```
In [37]: def load_models():
'''
    Replace '..path/' by the path of the saved models.
'''

    # Load the vectoriser.
    file = open('..path/vectoriser-ngram-(1,2).pickle', 'rb')
    vectoriser = pickle.load(file)
    file.close()
    # Load the LR Model.
    file = open('..path/Sentiment-LRv1.pickle', 'rb')
    LRmodel = pickle.load(file)
    file.close()

    return vectoriser, LRmodel

def predict(vectoriser, model, text):
    # Predict the sentiment
    textdata = vectorizer.transform(preprocess(text))
    sentiment = model.predict(textdata)

    # Make a List of text with sentiment.
    data = []
    for text, pred in zip(text, sentiment):
        data.append((text,pred))

    # Convert the List into a Pandas DataFrame.
    df = pd.DataFrame(data, columns = ['text','sentiment'])
    df = df.replace([0,1], ["Negative","Positive"])
    return df

if __name__=="__main__":
    # Loading the models.
    #vectoriser, LRmodel = load_models()

    # Text to classify should be in a List.
    text = ["I hate our president",
            "I Love you.",
            "Yes! We can win"]

    df = predict(vectorizer, LRmodel, text)
    print(df.head())
```

	text	sentiment
0	I hate our president	Negative
1	I Love you.	Positive
2	Yes! We can win	Positive

In []:

